

Incident Post-Mortem: US-East Database Outage

INCIDENT POST-MORTEM REPORT

Incident ID: INC-88219 Date of Incident: August 14, 2026 Lead Investigator: Marcus Thorne, Site Reliability Engineering

Overview: At 08:14 UTC, our primary database cluster in the US-East-2 region experienced a catastrophic failure, resulting in a 42-minute complete service outage for the core application suite. Time to Detect (TTD) was 3 minutes, and Time to Resolve (TTR) was 39 minutes.

Root Cause Analysis: The outage was triggered by a rogue automated maintenance script. The script was designed to purge orphaned temporary tables but contained a regex error that matched against active production partition keys. Upon execution, the script dropped three critical indexing tables, causing the primary PostgreSQL nodes to enter a fail-safe read-only state.

The automated failover mechanism to the US-West secondary cluster failed to initialize due to a misconfigured heartbeat timeout setting introduced in the previous week's infrastructure update. Consequently, application servers were unable to establish database connections, throwing HTTP 500 errors to end-users.

Resolution: The SRE team manually intervened at 08:22 UTC. We terminated the rogue script, restored the dropped index tables from the 08:00 UTC continuous snapshot, and forced a manual DNS failover to the secondary cluster to restore read/write capabilities while the primary cluster was rebuilt.

Action Items & Prevention: 1. Update CI/CD pipelines to require mandatory peer review for all database schema manipulation scripts. 2. Adjust the heartbeat timeout threshold on the failover routing policy from 15 seconds to 30 seconds to prevent false-negative health checks. 3. Implement a 'dry-run' execution environment for all maintenance cron jobs before they are allowed to interact with the production schema.